

## Tentative solution to ReEx Module 5: Resampling

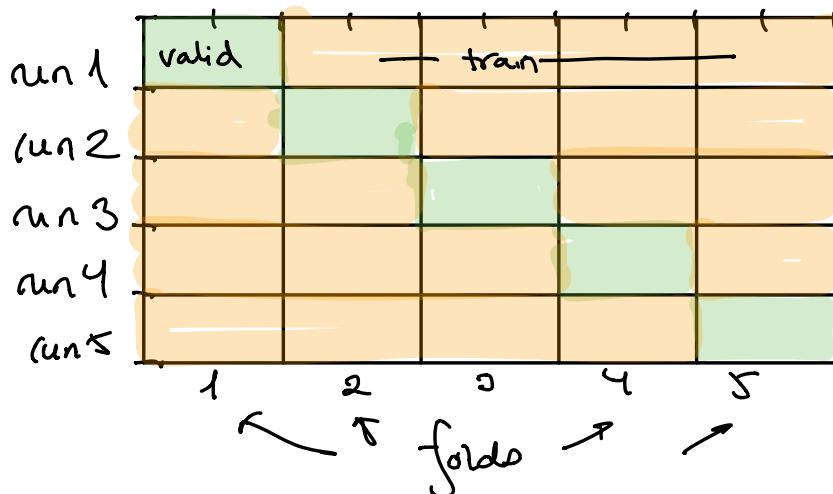
### Cross-validation

Q) Explain how  $k$ -fold cross-validation is implemented.

Drawing:  $k = 5$  (for simplicity)

First - shuffle the data: indices  $[1, 2, 17, \dots, 1, 28]$

Then partition into  $k$  groups = folds.



in general

$$\frac{n!}{k! \left[ \binom{n}{k} \right]^k}$$

possible ways to do this  
(multinomial with a twist)  
see below

train =

validate =

In run 1 fold 1 is kept aside and folds 2-5 are used to train the method (maybe many times, once for every model complexity). Then error is calculated on the validation fold.

Repeat  $k$  times:  $CV_n = \frac{1}{n} \sum_{j=1}^k MSE_j \cdot \eta_j$

$$nse_j = \frac{1}{n_j} \sum_{i \in G} (y_i - \hat{y}_i)^2 \text{ for } nse$$

↑  
 obs in validation fold  
 prediction of  $x_i$  in validation fold using fitted model from folds "i-j"  
 not j.

other loss functions may be 0/1 Loss.

Regression: find the <sup>optional</sup> number of neighbors in kNN-regression.

Classification: choose between QDA or LDA in classification

n obj. teile inn i k grupper und m in h ch,  $n = km$

(aut. perm. av de n) =  $n!$

(aut. perm. av de n) = (aut. måten å dele inn i k grupper)

$\cdot$  (aut. perm. av gruppene) =  $k!$   $\cdot$  (aut. måten å permutere innen grupper)

$$= B k! (m!)^k$$

$$B = \frac{n!}{k! (m!)^k} = \binom{n}{k} \frac{(n-k)!}{(m!)^k}$$

2) Advantages <sup>A</sup> & disadvantages <sup>D</sup> of k-fold CV relative to

a) the validation set

D: computational complexity

A: bias = generally larger sample size for each  $k$ -fold than validation set, which means "more data  $\rightarrow$  better fit" and therefore not overestimate the test set error

Bias? : compared to using the full data set for model fit.

A: different validation sets may give very different test error, so the results are unreliable - much more than for k-fold.

b) LOOCV = no randomness in splits!

A: less computational efforts for  $k$  than  $n$ , unless nice formula as for multiple linear regression.

A: less bias - in the sense that LOOCV uses a larger set to fit to data ( $n-1$  obs) which gives a less biased version of the test set error.

D: [According to our textbook]  
higher variance: we are averaging the output from  $n$  fitted models that are trained on nearly the same data  $\Rightarrow$  correlated positively.<sup>3</sup>

This happens to a less degree with k-fold, since the k models share different data and are thus less variable.

$$CV_n = \left( (y_1 - \hat{y}_1)^2 + (y_2 - \hat{y}_2)^2 + \dots + (y_n - \hat{y}_n)^2 \right) / n$$

$$\text{Var}(CV_n) = \frac{\text{sum variances of terms}}{\text{2 · covariances of terms}}$$



this part tend to be larger for LOOCV than k-fold  
when correlation higher between models

c) choice of k in k-fold. We just know that

$k=n = \text{LOOCV}$  (small bias - high variance) ↗ of estimator for test  
+ generally high comp. demand ↗ set or error

$k=2$  (larger bias - lower variance)  
+ less comp. challenging

simulations

and empirical research has found  $k=5$  or  $k=10$  to be good choices!

3) Case (as in R-code) : classification set up with two classes

-  $n=50$  observations of  $p=5000$  predictors

a) choose to use only  $d=25$  predictors, but choose the top  $d$  from absolute correlation coeff between the  $p$  preds. and the class label.

b) then use logistic regression with the  $d$  predictors.

⇒ How to do CV? On  $\underbrace{a+b}$  or only on  $\underbrace{b}$ ?

right    wrong

wrong: if only  $b$ , then all data used to find the predictors → gives overoptimistic result  
(misclassification rate of 0% can be found)

right: both  $a+b \Rightarrow$  all is good

See R-code in problem and run to see what the misclassification rates are.

## Bootstrapping

1)

a)  $P(\text{draw } x_i) = \frac{1}{n}, P(\text{not draw } x_i) = 1 - \frac{1}{n}$

b)  $P(\text{not any } x_i's) = (1 - \frac{1}{n})^n$

$P(\text{at least one } x_i) = 1 - (1 - \frac{1}{n})^n$

c)  $P(x_i \text{ in boot sample}) = 1 - (1 - \frac{1}{n})^n \approx 1 - \exp(-1)$   
 $= 0.632$

d) R-code to check result and see how fast  
 $1 - (1 - \frac{1}{n})^n \rightarrow 0.632$  (in n).

2) Bootstrap to estimate  $SD(\hat{\beta})$ :

for ( $b$  in 1:B) {

Draw with replacement from data to get

$(\bar{X}, \bar{Y})_b^*$ : bootstrap sample b,  $b=1, \dots, B$ .

fit  $\bar{Y} = \bar{X}\hat{\beta} + \varepsilon$  and keep  $\hat{\beta}_b$

}

Calculate  $SD(\hat{\beta}) = \sqrt{\frac{1}{B-1} \sum_{b=1}^B \left( \hat{\beta}_b - \frac{1}{B} \sum_{b=1}^B \hat{\beta}_b \right)^2}$

Why do we want to do this - when we really know that  $\hat{SD}(\hat{\beta}) = \hat{\sigma} \cdot \text{diag}((K\mathbf{X})^{-1})$ ?

And we might also do  $\hat{\text{Cov}}(\hat{\beta})$  but then use

$$\hat{\text{Cov}}(\hat{\beta}) = \frac{1}{B-1} \sum_{b=1}^B (\hat{\beta}_b - \bar{\hat{\beta}}_b)(\hat{\beta}_b - \bar{\hat{\beta}}_b)^T$$

$$\bar{\hat{\beta}}_b = \frac{1}{B} \sum_{b=1}^B \hat{\beta}_b$$

3) is covered on page 195 of the ISL book